

3. Estimation and Testing II

3.2 Hypothesis Testing

Hypothesis testing or significance testing is a method for testing a claim or hypothesis about a parameter in a population, using data measured in a sample.

We break up significance testing into four steps:

Step 1: Setting Null and Alternative Hypotheses

The first step of hypothesis testing is to convert the research question into null and alternative hypotheses.

- We start with the null hypothesis (H_0). The null hypothesis is a claim of “no difference”.
- The opposing hypothesis is the alternative hypothesis (H_1 or H_a). The alternative hypothesis is a claim of “a difference in the population” and is the hypothesis the researcher often hopes to bolster.

It is important to keep in mind that the null and alternative hypotheses reference population values, and not observed statistics.

Step 2: Calculating the Test statistic

We calculate a test statistic from the data. There are different types of test statistics. This chapter later describes those different types. Large test statistics indicate data are far from expected, providing evidence against the null hypothesis and in favor of the alternative hypothesis.

Step 3: p Value and Decision

The test statistic is converted to a probability called a p-value. The p-value answers the question “If the null hypothesis were true, what is the probability of observing the current data or data that is more extreme?”

Small p-values provide evidence against the null hypothesis because they say the observed data are unlikely when the null hypothesis is true. We apply the following convention:

- When $p\text{-value} \leq \alpha$, then the null hypothesis should be rejected in favor of alternative at α % level of significance.

Step 4: Conclusion

With decision made using the p-value, the conclusion should be written in non-technical form.

One-tailed and Two-tailed Testing

The equality part of the hypotheses always appears in the null hypothesis. In general, a hypothesis test about the value of a population mean μ must take one of the following three forms (where μ_0 is the hypothesized value of the population mean):

$$H_0: \mu \geq \mu_0$$


$$H_a: \mu < \mu_0$$

$$H_0: \mu \leq \mu_0$$

$$H_a: \mu > \mu_0$$

$$H_0: \mu = \mu_0$$

$$H_a: \mu \neq \mu_0$$



Depending on whether a hypothesis is one-tailed or two-tailed, the way of calculating p-value differs.

Type I and Type II Errors

Because hypothesis tests are based on sample data, we must allow for the possibility of errors.

- A Type I error is rejecting H_0 when it is true.

The probability of making a Type I error when the null hypothesis is true as an equality is called the level of significance. Applications of hypothesis testing that only control the Type I error are often called significance tests.

- A Type II error is accepting H_0 when it is false.

It is difficult to control for the probability of making a Type II error. Statisticians avoid the risk of making a Type II error by using “do not reject H_0 ” and not “accept H_0 ”.

3.2.1 Testing Single Population Mean (One Sample Z-test)

The one sample z-test is used to test a single population mean against a given value. In addition, this test is used only when the population standard deviation σ is known from a prior source.

Here, the null and alternative hypothesis can take any of the 3 forms:

$$\begin{array}{l} H_0: \mu \geq \mu_0 \\ H_a: \mu < \mu_0 \end{array}$$

$$\begin{array}{l} H_0: \mu \leq \mu_0 \\ H_a: \mu > \mu_0 \end{array}$$

$$\begin{array}{l} H_0: \mu = \mu_0 \\ H_a: \mu \neq \mu_0 \end{array}$$

Let us define the test statistic z in terms of the sample mean, the sample size and the population standard deviation σ :

$$z = \frac{\bar{x} - \mu_0}{\sigma / \sqrt{n}}$$

Then the null hypothesis should be rejected if the p-value corresponding to the calculated z value is less than α , where α is the desired level of significance.

Example:

Suppose a manufacturer claims that the mean lifetime of a light bulb is at least 10,000 hours. In a sample of 30 light bulbs, it was found that they only last 9,900 hours on average. Assume the population standard deviation is 120 hours. At 5% significance level, can we reject the claim by the manufacturer?

The hypothesis of interest is:

$$H_0: \mu \geq 10000$$

$$H_a: \mu < 10000$$

```
> xbar <- 9900           # sample mean
> mu0 <-10000           # hypothesized value
> sigma <-120           # population standard deviation
> n <-30                # sample size
> z <-(xbar-mu0)/(sigma/sqrt(n))
> z                     # test statistic
[1] -4.564355
>
> pval <-pnorm(z)
> pval                  # lower tail p-value
[1] 2.505166e-06
```

As the p-value (2.505e-06) turns out to be less than 0.05, we reject the null hypothesis at 5% significance level. So the mean lifetime of a light bulb is not above 10,000 hours.

Example:

Suppose the mean weight of King Penguins found in an Antarctic colony last year was 15.4 kg. In a sample of 35 penguins same time this year in the same colony, the mean penguin weight is 14.6 kg. Assume the population standard deviation is 2.5 kg. At 0.05 level of significance, can we reject the hypothesis that the mean penguin weight does not differ from last year?

The hypothesis of interest is :

$$H_0: \mu = 15.4$$

$$H_a: \mu \neq 15.4$$

```
> xbar <-14.6           # sample mean
> mu0 <-15.4           # hypothesized value
> sigma <-2.5          # population standard deviation
> n <-35               # sample size
> z <-(xbar-mu0)/(sigma/sqrt(n))
> z                    # test statistic
[1] -1.893146
> pval <-2 * pnorm(z)
> pval                 # two-tailed p-value
[1] 0.05833852
```

Here the p-value should be doubled, as the hypothesis is a two tailed one.

Since the final p-value turns out to be greater than the .05, we do not reject the null hypothesis at 5% significant level.

3.2.2 Testing Single Population Mean (One Sample t test)

The prior section used a z statistic to test a sample mean against an expectation. The z statistic needed population standard deviation σ (without estimating it from the data) to calculate the test statistic.

To conduct a one-sample test when the population standard deviation is not known, we use a variant of the z statistic called the t statistic. The advantage of the t statistic is that it can use sample standard deviation s instead of population standard deviation σ .

The null and alternative hypotheses are identical to those used by the z test.

The null hypothesis is $H_0: \mu = \mu_0$

Alternatives are,

$$H_a: \mu \neq \mu_0 \text{ (two-sided)}$$

$$H_a: \mu > \mu_0 \text{ (one-sided to right)}$$

$$H_a: \mu < \mu_0 \text{ (one-sided to left)}$$

Let us define the test statistic t in terms of the sample mean, the sample size and the sample standard deviation s as:

$$t = \frac{\bar{x} - \mu_0}{s / \sqrt{n}}$$

Here, the p-value corresponding to the calculated t value would be taken from the Student t-distribution with $n - 1$ degrees of freedom. Once the p-value is found, then the null hypothesis should be rejected at α % level of significance, if the p-value $\leq \alpha$.

Example:

Suppose a manufacturer claims that the mean lifetime of a light bulb is at least 10,000 hours. In a sample of 30 light bulbs, it was found that they only last 9,900 hours on average. Assume the sample standard deviation is 125 hours. At 0.05 significant level, can we reject the claim by the manufacturer?

The hypothesis of interest is :

$$H_0: \mu \geq 10000$$

$$H_a: \mu < 10000$$

```
> xbar <- 9900          # sample mean
> mu0 <- 10000          # hypothesized value
> s <- 125              # sample standard deviation
> n <- 30               # sample size
> t <- (xbar - mu0) / (s / sqrt(n))
> t                     # test statistic
[1] -4.38178
> pval <- pt(t, df=n-1)
> pval                  # lower tail p-value
[1] 7.035026e-05
```

The lower tail p-value of the test statistic (7.035e-05) turns out to be less than 0.05. We reject the null hypothesis at 5% significance level. So the mean lifetime of a light bulb is not above 10,000 hours.

3.2.3 Testing Single Population Proportion

A z test can be used to test a proportion from a single population. To carry out the test usually the sample should include at least 10 successes and 10 failures (Some texts say that 5 successes and 5 failures are enough.) and the population size is at least 10 times as big as the sample size.

Here, the null and alternative hypothesis can take any of the 3 forms:

$$\begin{array}{lll} H_0: p \geq p_0 & H_0: p \leq p_0 & H_0: p = p_0 \\ H_a: p < p_0 & H_a: p > p_0 & H_a: p \neq p_0 \end{array}$$

Let us define the test statistic z in terms of the sample proportion, sample size and the hypothesized value:

$$z = \frac{\bar{p} - p_0}{\sqrt{p_0(1 - p_0)/n}}$$

p-value can be found using the z distribution as in one sample z-test and the null hypothesis should be rejected if the p-value $\leq \alpha$, where α is the given level of significance.

Example:

Suppose that 12% of apples harvested in an orchard last year was rotten. 30 out of 214 apples in a harvest sample this year turns out to be rotten. At 0.05 significance level, test the hypothesis that the proportion of rotten apples in harvest is more than 12% this year?

The hypothesis of interest is :

$$H_0: p \leq 0.12$$

$$H_a: p > 0.12$$

```
> pbar = 30/214          # sample proportion
> p0 = .12              # hypothesized value
> n = 214               # sample size
> z = (pbar-p0)/sqrt(p0*(1-p0)/n)
> z                     # test statistic
[1] 0.908751
> pval = pnorm(z, lower.tail=FALSE)
> pval                  # upper tail p-value
[1] 0.1817408
```

The upper tail p-value of the test statistic (0.18) turns out to be greater than 0.05. We do not reject the null hypothesis at 5% significance level. So the proportion of rotten apples in harvest is not more than 12% this year.

3.2.4 Testing the Difference between Two Independent Population Means : σ_1 and σ_2 unknown but equal

We compare two independent samples to determine if the means for each sample are statistically significantly different. Either the two samples come from the same population and the populations mean difference is zero, or the two samples come from different populations and the population mean difference is not zero. The two samples do not have to be the same size.

Here, the null and alternative hypothesis can take any of the 3 forms:

$$\begin{array}{lll} H_0: \mu_1 - \mu_2 \geq 0 & H_0: \mu_1 - \mu_2 \leq 0 & H_0: \mu_1 - \mu_2 = 0 \\ H_a: \mu_1 - \mu_2 < 0 & H_a: \mu_1 - \mu_2 > 0 & H_a: \mu_1 - \mu_2 \neq 0 \end{array}$$

Let us define the test statistic t in terms of the two sample means \bar{x}_1, \bar{x}_2 , the sample sizes n_1, n_2 and the sample standard deviations s_1^2, s_2^2 :

$$t = \frac{\bar{x}_1 - \bar{x}_2}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

Where S_p is called the pooled sample variance and can be calculated using the formula :

$$s_p = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}}$$

Here the degree of freedom for the t value is $n_1 + n_2 - 2$.

Here, the p-value would be taken from the Student t-distribution with $n_1 + n_2 - 2$ degrees of freedom. Then the null hypothesis should be rejected at α % level of significance, if the p-value $\leq \alpha$.

Example:

Consider the dataset ‘mtcars’. Assume that the population variances for manual transmission car mileages and automatic transmission car mileages are equal. Test the hypothesis that the mean gas mileages of manual and automatic transmission cars are the same. Use 5% level of significance.

The hypothesis of interest is:

$$H_0: \mu_a - \mu_m = 0$$

$$H_a: \mu_a - \mu_m \neq 0$$

```
> attach(mtcars)
> auto<-mpg[am==0]
> manual<-mpg[am==1]
> t.test(auto,manual,var.equal = T)

Two Sample t-test
```

```
data:  auto and manual
t = -4.1061, df = 30, p-value = 0.000285
alternative hypothesis: true difference in means is not
equal to 0
95 percent confidence interval:
 -10.84837  -3.64151
sample estimates:
mean of x mean of y
 17.14737  24.39231
```

Here the p-value of the test is 0.00028, which is <0.05 . So at 5% level of significance we reject the null hypothesis of equal means of two populations. Hence, the millages for auto transmission and manual transmission cars are not the same.

3.2.5 Testing the Difference between Two Population Means: Paired Samples

Paired data involves taking two measurements on the same subjects, called repeated sampling.

Think of studying the effectiveness of a diet plan. You would weigh yourself prior to starting the diet and again following some time on the diet. Depending on how much weight you lost, you would determine if the diet was effective. What you might be interested in is estimating the true difference of the original weight and the weight lost after a certain period.

This test is testing the null hypothesis that there are no differences between the means of the two related groups. If we can reject the null hypothesis, then there is a significant difference between the means of the two groups.

So the hypotheses of interest is,

$$\begin{array}{lll} H_0: \mu_d \geq 0 & H_0: \mu_d \leq 0 & H_0: \mu_d = 0 \\ H_a: \mu_d < 0 & H_a: \mu_d > 0 & H_a: \mu_d \neq 0 \end{array}$$

where μ_d is the population mean difference of the two groups.

Then the test statistic of interest is,

$$t = \frac{\bar{x}_d}{s_d / \sqrt{n}}$$

where \bar{x}_d is the mean of the differences in the two samples, s_d is the standard deviation of the differences in the two samples and n is the sample size (number of pairs). This test statistic has a student t-distribution with $n-1$ degree of freedom.

If the p-value corresponding to the calculated t value is $\leq \alpha$, then the null hypothesis should be rejected at α % level of significance.

Example:

In the built-in data set named ‘immer’, the barley yields in years 1931 and 1932 of the same field are recorded. Assuming that the data in ‘immer’ follows the normal distribution, test the hypothesis that the mean barley yield in year 1931 (Y1) is larger than the mean barley yield in year 1932 (Y2).

The hypothesis of interest is:

$$\begin{aligned} H_0: \mu_d &\leq 0 & d = 1931 \text{ yield} - 1932 \text{ yield} \\ H_a: \mu_d &> 0 \end{aligned}$$

```
> attach(immer)
> t.test(Y1,Y2,paired=T, alternative="greater")
```

Paired t-test

```
data:  Y1 and Y2
t = 3.324, df = 29, p-value = 0.001206
alternative hypothesis: true difference in means is
greater than 0
95 percent confidence interval:
 7.778893      Inf
sample estimates:
mean of the differences
      15.91333
```

The p-value of the test is 0.0012, which is <0.05 . So at 5% level of significance we reject the null hypothesis. Hence the yield in 1931 is higher than the yield in 1932.